


Databases & SQL

1

1

A Database Schema



1) Stephens, R.K. and Plew, R.R., 2001. *Database Design*. SAMS, Indianapolis, IN. (with slight changes by V.G.D.)

2

1

Table

- “A *table* is the primary unit of physical storage for data in a database.”¹
- Usually a database contains more than one table.

1) Stephens, R.K. and Plew. R.R., 2001. *Database Design*. SAMS, Indianapolis , IN.

3

Table


Name	Company	Phone Number	E-mail Address
Vedat Diker	CLIS/UMD	(301) 405 9814	vedat@umd.edu
Bugs Bunny	Acme, Inc.	(123) 555 9876	bugs@acme.com
Will E. Coyote	Acme, Inc.	(123) 555 9821	will@acme.com

4

2

A Database with Multiple Tables

[1]



1) Stephens, R.K. and Plew. R.R., 2001. *Database Design*. SAMS, Indianapolis , IN. (with slight changes by V.G.D.)

5

Table

Customers

Name	Company	Phone Number	E-mail Address
Vedat Diker	CLIS/UMD	(301) 405 9814	vedat@umd.edu
Bugs Bunny	Acme, Inc.	(123) 555 9876	bugs@acme.com
Will E. Coyote	Acme, Inc.	(123) 555 9821	will@acme.com

6

3

Field (Column)

Customers

Name	Company	Phone Number	E-mail Address
Vedat Diker	CLIS/UMD	(301) 405 9814	vedat@umd.edu
Bugs Bunny	Acme, Inc.	(123) 555 9876	bugs@acme.com
Will E. Coyote	Acme, Inc.	(123) 555 9821	will@acme.com

a field

7

Record (Row)

Customers

Name	Company	Phone Number	E-mail Address
Vedat Diker	CLIS/UMD	(301) 405 9814	vedat@umd.edu
Bugs Bunny	Acme, Inc.	(123) 555 9876	bugs@acme.com
Will E. Coyote	Acme, Inc.	(123) 555 9821	will@acme.com

a record

8

4

Primary Key

Customers

Customer ID	Name	Company	Phone Number	E-mail Address
6273	Vedat Diker	CLIS/UMD	(301) 405 9814	vedat@umd.edu
3245	Bugs Bunny	Acme, Inc.	(123) 555 9876	bugs@acme.com
1324	Will E. Coyote	Acme, Inc.	(123) 555 9821	will@acme.com

primary key field

Primary key is a unique identifier of records in a table.

Primary key values may be generated manually or automatically.

9

Primary Key

Roles (Performances)

Actor/Actress	Movie	Character Name
Keanu Reeves	Matrix	Neo
Laurence Fishburne	Matrix	Morpheus
Carrie-Anne Moss	Matrix	Trinity
Keanu Reeves	Sweet November	Nelson Moss
Charlize Theron	Sweet November	Sara Deever
Charlize Theron	Waking Up in Reno	Candy Kirkendall
Laurence Fishburne	Othello	Othello
Ted Lange	Othello	Othello

primary key fields

A primary key can consist of more than one field.

10

SQL Introduction

Standard language for querying and manipulating data

Structured Query Language

Many standards out there:

- ANSI SQL
- SQL92 (a.k.a. SQL2)
- SQL99 (a.k.a. SQL3)
- Vendors support various subsets of these
- What we discuss is common to all of them

11

11

SQL

- Data Definition Language (DDL)
 - Create/alter/delete tables and their attributes
 - Following lectures...
- Data Manipulation Language (DML)
 - Query one or more tables – discussed next !
 - Insert/delete/modify tuples in tables
- Transact-SQL
 - Idea: package a sequence of SQL statements → server
 - Won't discuss in class

12

12

Data in SQL

1. Atomic types, a.k.a. data types
2. Tables built from atomic types

Unlike XML, no nested tables, only flat tables are allowed!

- We will see later how to decompose complex structures into multiple flat tables

13

13

Data Types in SQL

- Characters:
 - CHAR(20) -- fixed length
 - VARCHAR(40) -- variable length
- Numbers:
 - BIGINT, INT, SMALLINT, TINYINT
 - REAL, FLOAT -- differ in precision
 - MONEY
- Times and dates:
 - DATE
 - DATETIME -- SQL Server
- Others... All are simple

14

14

Tables in SQL

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

Tuples or rows

15

15

Tables Explained

- A tuple = a record
 - Restriction: all attributes are of atomic type
- A table = a set of tuples
 - Like a list...
 - ...but it is unordered: no **first()**, no **next()**, no **last()**.

16

16

Tables Explained

- The *schema* of a table is the table name and its attributes:
`Product(PName, Price, Category, Manfacturer)`
- A *key* is an attribute whose values are unique; we underline a key

`Product(PName, Price, Category, Manfacturer)`

17

17

SQL Query

Basic form: (plus many many more bells and whistles)

```
SELECT attributes  
FROM relations (possibly multiple)  
WHERE conditions (selections)
```

18

18

Simple SQL Query

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT *
FROM Product
WHERE category='Gadgets'
```



“selection”

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks

19

19

Simple SQL Query

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT PName, Price, Manufacturer
FROM Product
WHERE Price > 100
```



“selection” and
“projection”

PName	Price	Manufacturer
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

20

20

A Notation for SQL Queries

Input Schema

Product(PName, Price, Category, Manufacturer)

```
SELECT PName, Price, Manufacturer  
FROM Product  
WHERE Price > 100
```



Answer(PName, Price, Manufacturer)

Output Schema

21

21

Selections

What goes in the WHERE clause:

- $x = y$, $x < y$, $x \leq y$, etc
 - For numbers, they have the usual meanings
 - For CHAR and VARCHAR: lexicographic ordering
 - Expected conversion between CHAR and VARCHAR
 - For dates and times, what you expect...
- Pattern matching on strings...

22

22

The **LIKE** operator

- **s LIKE p:** pattern matching on strings
- **p** may contain two special symbols:
 - **%** = any sequence of characters
 - **_** = any single character

Product(PName, Price, Category, Manufacturer)

Find all products whose name mentions ‘gizmo’:

```
SELECT *
FROM Products
WHERE PName LIKE '%gizmo%'
```

23

23

Eliminating Duplicates

```
SELECT DISTINCT category
FROM Product
```

Category
Gadgets
Photography
Household

Compare to:

```
SELECT category
FROM Product
```

Category
Gadgets
Gadgets
Photography
Household

24

24

Ordering the Results

```
SELECT pname, price, manufacturer  
FROM Product  
WHERE category='gizmo' AND price > 50  
ORDER BY price, pname
```

Ordering is ascending, unless you specify the DESC keyword.

Ties are broken by the second attribute on the ORDER BY list, etc.

25

25

Joins in SQL


- Connect two or more tables:

Product	PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks	
Powergizmo	\$29.99	Gadgets	GizmoWorks	
SingleTouch	\$149.99	Photography	Canon	
MultiTouch	\$203.99	Household	Hitachi	


What is
the connection
between
them ?

Company	Cname	StockPrice	Country
GizmoWorks	25	USA	
Canon	65	Japan	
Hitachi	15	Japan	

26



27



28

Joins

Product (pname, price, category, manufacturer)
Company (cname, stockPrice, country)

Find all products under \$200 manufactured in Japan;
return their names and prices.

```
SELECT pname, price  
FROM Product, Company  
WHERE manufacturer=cname AND country='Japan'  
      AND price <= 200
```

Join
between Product
and Company

29

29

Joins in SQL

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

Company

Cname	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

```
SELECT pname, price  
FROM Product, Company  
WHERE manufacturer=cname AND country='Japan'  
      AND price <= 200
```



PName	Price
SingleTouch	\$149.99

30

30

Joins

Product (pname, price, category, manufacturer)
Company (cname, stockPrice, country)

Find all countries that manufacture some product in the
'Gadgets' category.

```
SELECT country
FROM Product, Company
WHERE manufacturer=cname AND category='Gadgets'
```

31

31

Joins in SQL

Product

Name	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

Company

Cname	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

```
SELECT country
FROM Product, Company
WHERE manufacturer=cname AND category='Gadgets'
```

What is
the problem ?
What's the
solution ?



Country
??
??

32

32

Joins

Product (pname, price, category, manufacturer)

Purchase (buyer, seller, store, product)

Person(persname, phoneNumber, city)

Find names of people living in Seattle that bought some product in the ‘Gadgets’ category, and the names of the stores they bought such product from

```
SELECT DISTINCT persname, store  
FROM Person, Purchase, Product  
WHERE persname=buyer AND product = pname AND  
city='Seattle' AND category='Gadgets'
```

33

When are two tables related?

- You guess they are
- I tell you so
- Foreign keys are a method for schema designers to tell you so
 - A foreign key states that a column is a reference to the key of another table
 - ex: Product.manufacturer is foreign key of Company
 - Gives information and enforces constraint

34

34